

## 9\_api\_data

December 11, 2023

### 1 APIs: Exploring Guardian API Data

Guardian news data provides us a range of different types of variable that we can use to get an overall picture of our dataset, and perhaps even find a some interesting patterns along the way.

Below we look at a range of different options for examining Guardian Data. Whilst the text of the stories is obviously valuable data, we'll need more advanced text analysis methods for that. These methods allow us to get a good overall picture of our data and find general trends.

```
[ ]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[ ]: articles = pd.read_json('AI_articles.json')
articles.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2000 entries, 0 to 1999
Data columns (total 13 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   id                    2000 non-null   object
 1   type                  2000 non-null   object
 2   sectionId             2000 non-null   object
 3   sectionName           2000 non-null   object
 4   webPublicationDate     2000 non-null   object
 5   webTitle               2000 non-null   object
 6   webUrl                 2000 non-null   object
 7   apiUrl                 2000 non-null   object
 8   fields                 2000 non-null   object
 9   tags                  2000 non-null   object
10   isHosted              2000 non-null   bool
11   pillarId              1980 non-null   object
12   pillarName            1980 non-null   object
dtypes: bool(1), object(12)
memory usage: 205.1+ KB
```

## 1.1 Prepping the Data

First we get the data ready for analysis.

### 1.1.1 Transforming the date column

First we need to ensure that our data is clean and that the `webPublicationDate` is properly formatted as a `datetime`.

```
[ ]: articles.head()
```

```
[ ]:
      id      type  sectionId \
0  technology/2023/oct/31/educators-teachers-ai-l...  article  technology
1  technology/ng-interactive/2023/oct/25/a-day-in...  interactive  technology
2  technology/2023/oct/24/alphabet-q3-earnings-go...  article  technology
3  stage/2023/sep/19/anthropology-review-hampstea...  article      stage
4  film/2023/aug/20/tim-review-clunky-ai-paranoia...  article      film

      sectionName  webPublicationDate \
0  Technology  2023-10-31T10:00:39Z
1  Technology  2023-10-25T13:38:11Z
2  Technology  2023-10-24T22:07:37Z
3      Stage  2023-09-19T12:02:55Z
4      Film  2023-08-20T10:30:44Z

      webTitle \
0  'Is this an appropriate use of AI or not?': te...
1      A day in the life of AI
2  Google Cloud revenue misses expectations despi...
3  Anthropology review - clever AI missing-person...
4      TIM review - clunky AI paranoia thriller

      webUrl \
0  https://www.theguardian.com/technology/2023/oc...
1  https://www.theguardian.com/technology/ng-inte...
2  https://www.theguardian.com/technology/2023/oc...
3  https://www.theguardian.com/stage/2023/sep/19/...
4  https://www.theguardian.com/film/2023/aug/20/t...

      apiUrl \
0  https://content.guardianapis.com/technology/20...
1  https://content.guardianapis.com/technology/ng...
2  https://content.guardianapis.com/technology/20...
3  https://content.guardianapis.com/stage/2023/se...
4  https://content.guardianapis.com/film/2023/aug...

      fields \
0  {'byline': 'Johana Bhuiyan', 'body': '<p>In <a...
```

```

1 {'byline': 'Hannah Devlin Science Corresponden...
2 {'byline': 'Kari Paul', 'body': '<p>Google is ...
3 {'byline': 'Mark Lawson', 'body': '<p>While sc...
4 {'byline': 'Wendy Ide', 'body': '<p>This styli...

```

```

                                tags  isHosted  pillarId  \
0  [{'id': 'technology/technology', 'type': 'keyw...  False  pillar/news
1  [{'id': 'technology/artificialintelligenceai',...  False  pillar/news
2  [{'id': 'technology/alphabet', 'type': 'keywor...  False  pillar/news
3  [{'id': 'stage/stage', 'type': 'keyword', 'sec...  False  pillar/arts
4  [{'id': 'film/thriller', 'type': 'keyword', 's...  False  pillar/arts

```

```

pillarName
0      News
1      News
2      News
3      Arts
4      Arts

```

```
[ ]: articles['webPublicationDate'] = pd.to_datetime(articles['webPublicationDate'])
```

### 1.1.2 Unpacking the Fields column

The content of the `fields` column is determined when we collect our API data, by what we passed to `show-fields` in our query parameters. However what is returned is a dictionary of information. Ideally we want to expand these dictionaries out and create additional columns for each field (byline, body and wordcount).

We'll mainly be using `wordcount` but the process will unpack all fields.

```
[ ]: articles.loc[0, 'fields']
```

```
[ ]: {'byline': 'Johana Bhuiyan',
      'body': '<p>In <a
href="https://www.theguardian.com/technology/2023/oct/12/chatgpt-uses-writing-
recipes-one-year">the year since OpenAI released ChatGPT</a>, high school
teacher Vicki Davis has been rethinking every single assignment she gives her
students. Davis, a computer science teacher at Sherwood Christian Academy in
Georgia, was well-positioned to be an early adopter of the technology. She's
also the IT director at the school and helped put together an AI policy in
March: the school opted to allow the use of AI tools for specific projects so
long as students discussed it with their teachers and cited the tool. In Davis's
mind, there were good and bad uses of AI, and ignoring its growing popularity
was not going to help students unlock the productive uses or understand its
dangers.</p> <aside class="element element-rich-link element--thumbnail"> <p>
<span>Related: </span><a
href="https://www.theguardian.com/technology/2023/oct/26/ai-artificial-
intelligence-investment-boom">Humanity at risk from AI 'race to the bottom',

```

says tech expert

"It's actually changed how I design my projects because there are some times I want my students to use AI, and then there are times I don't want them to," Davis said. "What am I trying to teach here? Is this an appropriate use of AI or not?"

Like teachers across the US and UK, Davis, who also runs the education blog Cool Cat Teacher, spent the summer thinking through what the release of a technology could mean for her.

Generative AI can produce images of the pope in a bomber jacket and answer nearly any math problem, so what could it do for students? Educators like her played with the tools and tried to understand how they work, what the utility could be - for teachers and students alike - and, perhaps most pressingly, how the software could be misused. Some took drastic measures, going so far as to abandon homework assignments as long as the technology was accessible.

"It feels like we're in some sort of lab experimenting with our kids because it's changing so rapidly," Davis said. "If you had asked me about any of this last fall, I couldn't have told you any of it because ChatGPT didn't exist."

In Davis's senior level class, she prohibited the use of chatbots to code because until recently the College Board, which administers standardized tests like the SAT, didn't permit AI assistance for programming. (This was [recently changed](https://apcentral.collegeboard.org/exam-administration-ordering-scores/administering-exams/preparing-for-exam-day/exam-security/artificial-intelligence-tools) to allow for the use of generative AI as a supplemental tool.) But she has changed an annual project she assigns to incorporate AI into the process. Davis usually asks students to research current models of laptops and evaluate which would be the best fit based on where they want to go to college and what they want to study. Now she asks students to feed the research they have done on their computer options into ChatGPT and ask for a recommendation based on their chosen major and college. The students are then tasked with evaluating ChatGPT's recommendation. Her goal is to show students how they can use their own knowledge and research on a topic to help them better supervise AI.

If you had asked me about any of this last fall, I couldn't have told you any of it because ChatGPT didn't exist

Vicki Davis

Teachers who spoke to the Guardian say their primary concern is helping students begin to use AI without enabling cheating. Looming over their futuristic lessons is a fear that an overreliance on these new tools could exacerbate the loss of learning many students suffered during the pandemic. Students had only returned to in-person instruction after two remote years when OpenAI launched ChatGPT, and many were still struggling with the huge hit to their ability to learn or engage in school at all.

"There's so much trauma, and AI can't help me with that," said one Maryland high school teacher, Kevin Shindel.

\*\*\*

After a summer spent experimenting with AI, there's little consensus among teachers on how to address its use in schools. Many educators in a nearly 370,000-person Facebook group called "ChatGPT for teachers" argue the widespread use of AI chatbots is inevitable and eagerly discuss the best ways to use these tools to make their jobs more efficient and help their students learn. Other teachers the Guardian spoke to suggested student use of the tools be banned until they learn

more about the technology behind it.

Still, others have focused largely on mitigating any AI-aided cheating; some have stopped assigning homework entirely, opting instead to have their students do supervised work in class. Some teachers have even required students to take handwritten exams or write the first drafts of essays by hand in class to ensure they are coming up with the ideas themselves.

But all those the Guardian spoke to agree: regardless of where you land on its use, teachers everywhere are grappling with how to stay on top of constantly evolving generative AI tools.

Shindel, a government teacher at a 3,300-student high school in Maryland, has been teaching his students about how AI impacts government and policy for 15 years, but he wasn't prepared for how quickly people would adopt ChatGPT. He spent the summer learning about and experimenting with various chatbots, and in July presented his findings to the school board in a 38-slide presentation titled **"The promise and peril of ChatGPT in today's classroom"**.

**Related:** [A day in the life of AI](https://www.theguardian.com/technology/ng-interactive/2023/oct/25/a-day-in-the-life-of-ai)

Shindel gave those in attendance ChatGPT-led activities to experiment with and posed questions about the ethics of its use ("What would a code of ethics for data usage and protection look like?"). Ultimately, he urged the school board to come up with a district-wide policy.

"Teachers shouldn't be responsible for developing classroom policies alone," Shindel said. "There needs to be some kind of concerted, systemic effort."

Shindel doesn't believe teachers and policymakers know enough about how chatbots collect student's personal information - or how to prevent cheating - to allow students to use it. He also worries the tools could exacerbate the lack of student engagement caused by remote learning. Students and teachers are still reeling from the impacts of the pandemic, Shindel said. A recent Harvard graduate school of education [study](https://www.gse.harvard.edu/ideas/news/23/05/new-data-show-how-pandemic-affected-learning-across-whole-communities#:~:text=Living%20in%20a%20community%20where,of%20a%20year%20in%20reading.>study) concluded the average public school student between third and eighth grade was half a year behind in math and reading and that nearly all students failed to recover the learning lost after returning to in-person instruction. A 2021 [review of 10 studies](https://www.gov.uk/government/publications/learning-during-the-pandemic/learning-during-the-pandemic-quantifying-lost-time--2) on pandemic learning loss published by the UK's Department for Education found that "disadvantaged primary school students were disproportionately behind expectations", with many students 50% further behind.

"I have a couple classes that are almost completely silent. Students don't interact with each other or answer any questions," Shindel said.

Though they may be in the minority, other schools have made progress establishing AI policies. Little Falls high school in Minnesota decided to ban the use of AI tools entirely in an addendum to the school-wide cheating policy. Davis's class policy allows certain tools to be used but requires students to seek permission and review the links the AI cites as sources. Kimberly Van Orman, a University of Georgia philosophy

professor who is currently teaching a course on the ethics of AI, says she is focusing on transparency. Van Orman requires her students to include the prompt they entered into a chatbot and the response in any assignment they use it for to ensure they don't "use it in a way that takes the place of learning".

"If you're trying to understand a concept from the book and you want to kind of talk it over with ChatGPT, that would be fine," Van Orman said. "Consulting it on your homework problem would not be fine."

\*\*\*\*

Dozens of AI apps targeting students have cropped up in the past few years. Photomath, for instance, predates the current versions of ChatGPT and pitches itself as the No 1 app for math learning. Users can upload a picture of a math problem or equation, and the app will give them the answer with explanations. But several teachers said students began using it during the pandemic to cheat or, at the very least, replace the "productive struggle" that results in learning. Inevitably, students who relied on Photomath during the pandemic struggled when they returned to the classroom, several teachers said.

aside class="element element-pullquote element--supporting">
 <blockquote>
 <p>There's so much trauma, and AI can't help me with that</p>
 </blockquote>
 </aside>
 <p>But there are also tools being built to refuse to just give students the answer. Khanmigo, an AI tutor being piloted by educational non-profit Khan Academy, is trained instead to ask questions that nudge students to better understand the material. When the Guardian asked Khanmigo a basic programming question (implement a cache with expirations in Javascript), the chatbot responded: "I can't provide direct answers or solutions to coding problems." When the Guardian was asked to solve for z in the equation "3z = 15" and repeatedly responded with "I don't know", the AI tutor kept providing guidance on how to solve it until it finally provided four multiple-choice options. Khanmigo was quicker to provide the right answer when the Guardian responded with an incorrect answer twice. ChatGPT, on the other hand, immediately provided the solution in both cases.

Sal Khan, the founder of Khan Academy, says the organization spent thousands of hours training the system, which is powered by ChatGPT-4, to understand that it's not supposed to do people's work for them. "We said stuff like: 'You're a Socratic tutor, you are here to make the students actively learn, not just passively,'" he said.

Though it's still in an experimental phase, these training processes are what distinguishes an AI tutor from an AI cheating tool, Khan argues.

"This time next year, you're going to have 50 [companies] who say that they have an AI tutor," Khan said. "But probably 90% of them are going to be somewhat shady and they're just going slap a little bit of a layer on top of ChatGPT-3.5. They're going to be mainly cheating tools, and not good ones."

'wordcount': '1585'}

.json\_normalize happens to do this kind of job for us, but it will create an entirely new dataframe from the results.

```
[ ]: articles_field_data = pd.json_normalize(articles['fields'])
articles_field_data.head()
```

```
[ ]:                                     byline \
0                                     Johana Bhuiyan
1 Hannah Devlin Science Correspondent, Rich Cous...
2                                     Kari Paul
3                                     Mark Lawson
4                                     Wendy Ide

                                     body wordcount
0 <p>In <a href="https://www.theguardian.com/tec... 1585
1 <figure class="element element-atom element--i... 1741
2 <p>Google is doing well, but not well enough f... 554
3 <p>While screenwriters strike, partly over the... 410
4 <p>This stylishly icy-looking thriller sounds ... 86
```

Having produced our dataframe of field data we just need to merge the `articles` dataframe and the new one together, matching up the indexes. When merging dataframes, left literally refers to the dataframe on the left of the operation, and right to the one most towards the right.

```
left.merge(right)
```

```
[ ]: articles = articles.merge(articles_field_data, left_index=True,
    ↪right_index=True)
articles.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2000 entries, 0 to 1999
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    2000 non-null  object
1   type                  2000 non-null  object
2   sectionId             2000 non-null  object
3   sectionName           2000 non-null  object
4   webPublicationDate     2000 non-null  datetime64[ns, UTC]
5   webTitle              2000 non-null  object
6   webUrl                2000 non-null  object
7   apiUrl                2000 non-null  object
8   fields                2000 non-null  object
9   tags                  2000 non-null  object
10  isHosted              2000 non-null  bool
11  pillarId              1980 non-null  object
12  pillarName            1980 non-null  object
13  byline                1944 non-null  object
14  body                  2000 non-null  object
15  wordcount             2000 non-null  object
dtypes: bool(1), datetime64[ns, UTC](1), object(14)
memory usage: 316.5+ KB
```

### 1.1.3 Converting wordcount to numeric

Wordcount has been stored as a string. We can rectify that by using `.to_numeric`

```
[ ]: articles['wordcount'] = pd.to_numeric(articles['wordcount'])
```

## 1.2 Data Counts over time

A key question of a dataset about the news, would when this news took place. Equally, we may also be interested in trends over time. Depending on your query, it may be interesting to see if there were changing publication rates related to your topic of interest.

A simple `.describe` on the date column will tell us a little about the spread of the dates.

```
[ ]: articles['webPublicationDate'].describe(datetime_is_numeric=True)
```

```
[ ]: count                2000
     mean    2023-03-31 20:56:41.787000064+00:00
     min           2020-08-21 14:20:31+00:00
     25%    2023-01-31 11:58:43.750000128+00:00
     50%           2023-05-25 13:13:57+00:00
     75%    2023-08-15 11:19:21.750000128+00:00
     max           2023-11-08 15:05:42+00:00
     Name: webPublicationDate, dtype: object
```

If we want to see trends, we can group our rows by publication period such as by Day, Month or Year. To do this we make a special time grouping object, and then group our data using it. We count the number of articles in each group and then plot them as a line plot.

```
[ ]: time_grouper = pd.Grouper(key='webPublicationDate', freq='M')
     count_over_time = articles[['webPublicationDate', 'id']].groupby(time_grouper).
     ↪count().reset_index()
     count_over_time.head()
```

```
[ ]:      webPublicationDate  id
0 2020-08-31 00:00:00+00:00    1
1 2020-09-30 00:00:00+00:00    0
2 2020-10-31 00:00:00+00:00    0
3 2020-11-30 00:00:00+00:00    2
4 2020-12-31 00:00:00+00:00    7
```


Time series are a little trickier to plot and Seaborn doesn't have a built in convenience method for it. However we can use the `sns.lineplot` method to manually create one, and make some adjustments to size and label positioning manually.

```
[ ]: plt.figure(figsize=(20,5))
     plot = sns.lineplot(data=count_over_time, x='webPublicationDate', y='id')

     plot.tick_params(axis='x', labelrotation=45)
     plot.set(title='AI story Frequency over time', xlabel='Month', ylabel='Freq')
```



```
sns.despine()
plt.show()
```



### 1.3 Optional: Filtering by a Date Range

Using our timeseries plot we might decide to filter our data so we only work with a specific range period.

```
[ ]: date_filter = articles['webPublicationDate'] >= 'January 2022'
articles = articles[date_filter]
```

```
[ ]: articles['webPublicationDate'].describe()
```

<ipython-input-13-a6523f342fc1>:1: FutureWarning: Treating datetime data as categorical rather than numeric in `.describe` is deprecated and will be removed in a future version of pandas. Specify `datetime\_is\_numeric=True` to silence this warning and adopt the future behavior now.

```
articles['webPublicationDate'].describe()
```

```
[ ]: count                1887
unique                1863
top    2023-03-26 09:00:14+00:00
freq                      3
first    2022-01-09 09:00:17+00:00
last     2023-11-08 15:05:42+00:00
Name: webPublicationDate, dtype: object
```

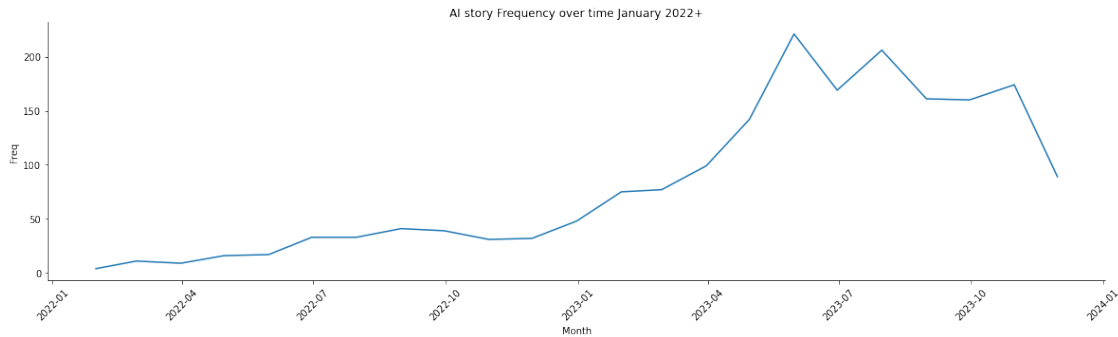
```
[ ]: time_grouper = pd.Grouper(key='webPublicationDate', freq='M')
count_over_time = articles[['webPublicationDate', 'id']].groupby(time_grouper).
    ↪count().reset_index()

plt.figure(figsize=(20,5))
plot = sns.lineplot(data=count_over_time, x='webPublicationDate', y='id')
plot.tick_params(axis='x', labelrotation=45)
```

```

plot.set(title='AI story Frequency over time January 2022+', xlabel='Month',
         ylabel='Freq')
sns.despine()
plt.show()

```



## 1.4 Appropriate Pillars?

The Guardian has a number of major sections they refer to as Pillars. We can examine the distribution of our articles across these major categories.

```

[ ]: pillar_counts = articles['pillarName'].value_counts()
pillar_counts

```

```

[ ]: News          1072
     Arts           479
     Opinion        201
     Lifestyle       73
     Sport           42
     Name: pillarName, dtype: int64

```

```

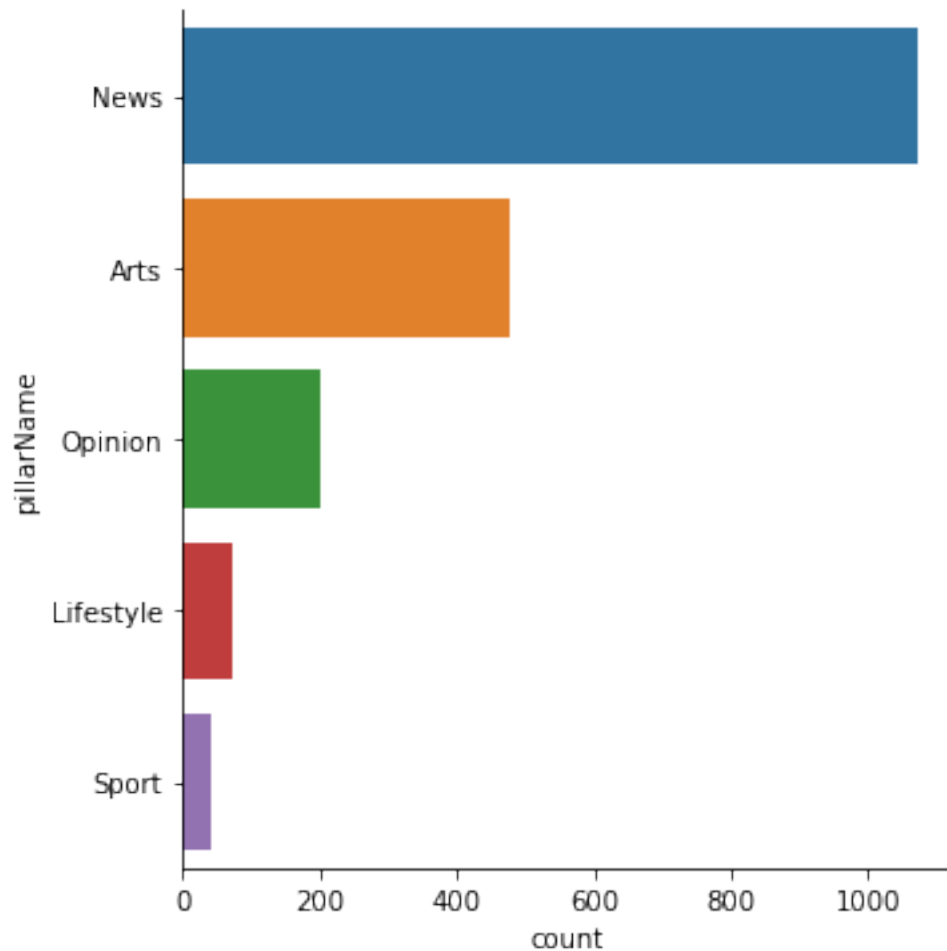
[ ]: sns.catplot(data=articles, y='pillarName', kind='count', order=pillar_counts.
         index)

```

```

[ ]: <seaborn.axisgrid.FacetGrid at 0x1207b3250>

```



## 1.5 Optional: Filtering by Pillar

Depending on your search query and the type of question you have, it may be worth filtering out material in unsuitable pillars, or focusing on just one.

```
[ ]: chosen_pillars = ['News', 'Opinion']
pillar_filter = articles['pillarName'].isin(chosen_pillars)
articles = articles[pillar_filter]
```


After filtering we can re-run our counts to check the filtering was applied, and produce a new visualisation of we need it.

```
[ ]: new_pillar_counts = articles['pillarName'].value_counts()
new_pillar_counts
```

```
[ ]: News      1072
Opinion     201
Name: pillarName, dtype: int64
```

```
[ ]: sns.catplot(data=articles, y='pillarName', kind='count',  
↳ order=new_pillar_counts.index)
```

```
[ ]: <seaborn.axisgrid.FacetGrid at 0x11bd564f0>
```



## 1.6 Sections

Sections are the next form of categorisation. Sections give us a better sense of the overall topic of the stories.

```
[ ]: section_counts = articles['sectionName'].value_counts()  
section_counts
```

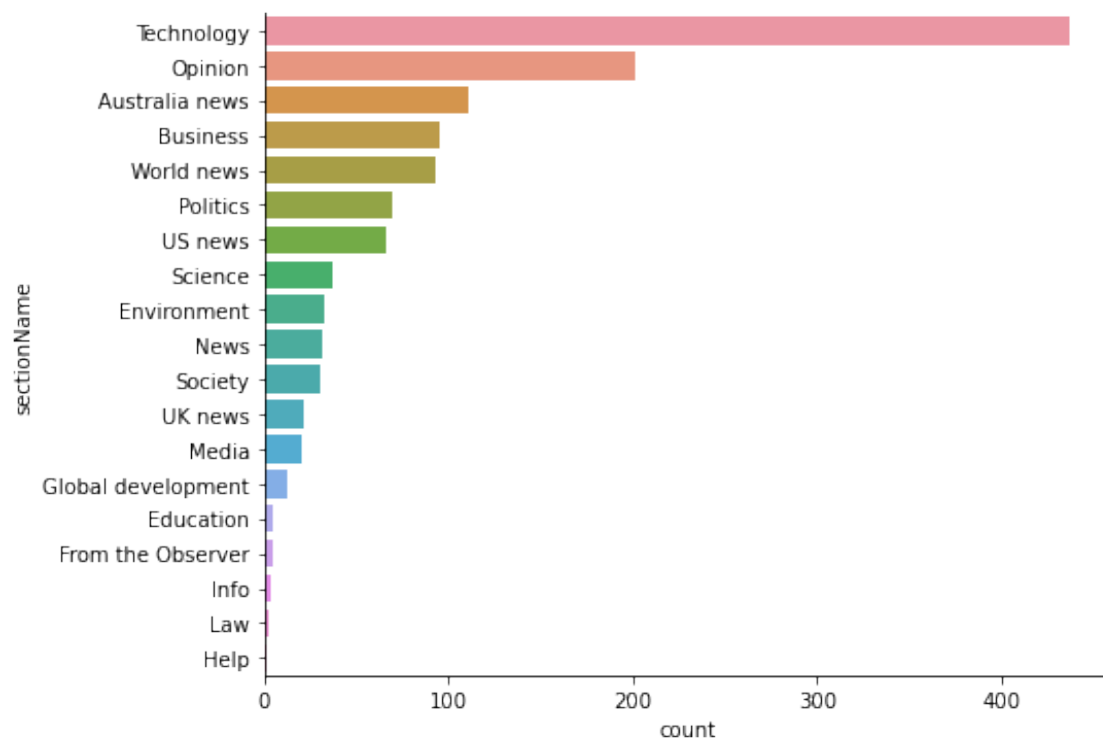
```
[ ]: Technology      437  
     Opinion        201  
     Australia news  111  
     Business        95
```

World news	93
Politics	69
US news	66
Science	37
Environment	33
News	31
Society	30
UK news	21
Media	20
Global development	12
Education	5
From the Observer	5
Info	4
Law	2
Help	1

Name: sectionName, dtype: int64

```
[ ]: sns.catplot(data=articles, y='sectionName', kind='count', aspect=1.5,
↳ order=section_counts.index)
```

```
[ ]: <seaborn.axisgrid.FacetGrid at 0x11be677c0>
```



Depending on how many sections are involved we may decide to keep only those above a certain threshold of presence in our dataset. This could be a top 10 or 20, or you could base it on some

sort of summary metric of the counts such as categories above the mean or median count.

```
[ ]: section_counts.describe()
```

```
[ ]: count      19.000000
     mean       67.000000
     std       102.893148
     min        1.000000
     25%        8.500000
     50%       31.000000
     75%       81.000000
     max      437.000000
     Name: sectionName, dtype: float64
```

```
[ ]: above_avg_sections = section_counts[section_counts > section_counts.median()].
     ↪index
     above_avg_sections
```

```
[ ]: Index(['Technology', 'Opinion', 'Australia news', 'Business', 'World news',
           'Politics', 'US news', 'Science', 'Environment'],
          dtype='object')
```

We'll just go with a top 10

```
[ ]: top_sections = section_counts.index[:10]
     top_sections
```

```
[ ]: Index(['Technology', 'Opinion', 'Australia news', 'Business', 'World news',
           'Politics', 'US news', 'Science', 'Environment', 'News'],
          dtype='object')
```

```
[ ]: articles = articles[articles['sectionName'].isin(top_sections)]
     articles.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1173 entries, 0 to 1997
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    1173 non-null   object
1   type                  1173 non-null   object
2   sectionId             1173 non-null   object
3   sectionName           1173 non-null   object
4   webPublicationDate     1173 non-null   datetime64[ns, UTC]
5   webTitle              1173 non-null   object
6   webUrl                1173 non-null   object
7   apiUrl               1173 non-null   object
8   fields                1173 non-null   object
9   tags                  1173 non-null   object
```

```

10  isHosted          1173 non-null    bool
11  pillarId          1173 non-null    object
12  pillarName        1173 non-null    object
13  byline            1138 non-null    object
14  body              1173 non-null    object
15  wordcount          1173 non-null    int64
dtypes: bool(1), datetime64[ns, UTC](1), int64(1), object(13)
memory usage: 147.8+ KB

```

### 1.6.1 Examining Section Content

We may find interesting sections in our dataset and wonder why they're there. We can iterate through the titles and URLs of the section we're interested in to get a better sense of why they've been included.

Below is a simple section filter but you could make it more complicated, such as limiting to after a time period, or in combination with a pillar classification for example.

N.B Below we use `.head` to limit the number of results for demonstration purposes, but during analysis there is no reason you cannot remove it and to view all the results.

```

[ ]: SECTION_OF_INTEREST = 'Australia news' # Just change this to switch sections

selected_data = articles[articles['sectionName'] == SECTION_OF_INTEREST].head(5)
for index, row in selected_data.iterrows():
    print(row['webTitle'])
    print(row['webUrl'])
    print('****')

```

Australian federal police using AI to analyse data obtained under surveillance warrants

<https://www.theguardian.com/australia-news/2023/sep/22/australian-federal-police-afp-using-ai-analyse-surveillance-warrants-data>

\*\*\*\*

Morning Mail: 'secretive' Israel defence exports, bombshell testimony in Trump trial, AI risks debated

<https://www.theguardian.com/australia-news/2023/oct/25/morning-mail-secret-israel-defence-exports-bombshell-testimony-in-trump-trial-ai-risks-debated>

\*\*\*\*

AI could 'turbo-charge fraud' and be monopolised by tech companies, Andrew Leigh warns

<https://www.theguardian.com/australia-news/2023/sep/20/ai-artificial-intelligence-warnings-dangers-andrew-leigh-mckell-institute>

\*\*\*\*

Morning Mail: authors fear huge AI copyright 'theft', scorching weekend for east, Michael Gambon dies

<https://www.theguardian.com/australia-news/2023/sep/29/morning-mail-authors-fear-huge-ai-copyright-theft-scorching-weekend-for-east-michael-gambon-dies>

\*\*\*\*

Democracies face ‘truth decay’ as AI blurs fact and fiction, warns head of Australia’s military

<https://www.theguardian.com/australia-news/2023/sep/14/democracies-face-truth-decay-as-ai-blurs-fact-and-fiction-warns-head-of-australias-military>

\*\*\*\*

## 1.7 Tags

Tags are the last categorisation and they give us even more nuance in exactly what each story is about. However they are a little trickier to deal with because each story can have more than one tag associated with it. This presents us more of a challenge but also an opportunity for analysis too.

If we look at the first item in our `tags` column, we can see that the value is actually a quite complex object. A list and then each item in the list is a dictionary. A lot of information is provided but for our purposes we just want the tag string, which is held under the `webTitle` key.

```
[ ]: articles['tags'].iloc[0]
```

```
[ ]: [{ 'id': 'technology/technology',
      'type': 'keyword',
      'sectionId': 'technology',
      'sectionName': 'Technology',
      'webTitle': 'Technology',
      'webUrl': 'https://www.theguardian.com/technology/technology',
      'apiUrl': 'https://content.guardianapis.com/technology/technology',
      'references': []},
     { 'id': 'technology/artificialintelligenceai',
      'type': 'keyword',
      'sectionId': 'technology',
      'sectionName': 'Technology',
      'webTitle': 'Artificial intelligence (AI)',
      'webUrl': 'https://www.theguardian.com/technology/artificialintelligenceai',
      'apiUrl': 'https://content.guardianapis.com/technology/artificialintelligenceai',
      'references': []},
     { 'id': 'education/education',
      'type': 'keyword',
      'sectionId': 'education',
      'sectionName': 'Education',
      'webTitle': 'Education',
      'webUrl': 'https://www.theguardian.com/education/education',
      'apiUrl': 'https://content.guardianapis.com/education/education',
      'references': []}]
```

As each story could have multiple tags we’re going to create a version of the `articles` dataframe where each row represents a single tag, and other story information like title, wordcount etc are duplicated. Pandas will keep track of which rows all refer to the same story using the index.



```
[ ]: tag_per_line = articles.explode('tags')
tag_per_line.head(10)
```

```
[ ]:
id      type  sectionId \
0  technology/2023/oct/31/educators-teachers-ai-l...  article  technology
0  technology/2023/oct/31/educators-teachers-ai-l...  article  technology
0  technology/2023/oct/31/educators-teachers-ai-l...  article  technology
1  technology/ng-interactive/2023/oct/25/a-day-in...  interactive  technology
1  technology/ng-interactive/2023/oct/25/a-day-in...  interactive  technology
1  technology/ng-interactive/2023/oct/25/a-day-in...  interactive  technology
1  technology/ng-interactive/2023/oct/25/a-day-in...  interactive  technology
2  technology/2023/oct/24/alphabet-q3-earnings-go...  article  technology
2  technology/2023/oct/24/alphabet-q3-earnings-go...  article  technology
2  technology/2023/oct/24/alphabet-q3-earnings-go...  article  technology
```

```
sectionName  webPublicationDate \
0  Technology 2023-10-31 10:00:39+00:00
0  Technology 2023-10-31 10:00:39+00:00
0  Technology 2023-10-31 10:00:39+00:00
1  Technology 2023-10-25 13:38:11+00:00
1  Technology 2023-10-25 13:38:11+00:00
1  Technology 2023-10-25 13:38:11+00:00
1  Technology 2023-10-25 13:38:11+00:00
2  Technology 2023-10-24 22:07:37+00:00
2  Technology 2023-10-24 22:07:37+00:00
2  Technology 2023-10-24 22:07:37+00:00
```

```
webTitle \
0  'Is this an appropriate use of AI or not?': te...
0  'Is this an appropriate use of AI or not?': te...
0  'Is this an appropriate use of AI or not?': te...
1  A day in the life of AI
1  A day in the life of AI
1  A day in the life of AI
1  A day in the life of AI
2  Google Cloud revenue misses expectations despi...
2  Google Cloud revenue misses expectations despi...
2  Google Cloud revenue misses expectations despi...
```

```
webUrl \
0  https://www.theguardian.com/technology/2023/oc...
0  https://www.theguardian.com/technology/2023/oc...
0  https://www.theguardian.com/technology/2023/oc...
1  https://www.theguardian.com/technology/ng-inte...
1  https://www.theguardian.com/technology/ng-inte...
1  https://www.theguardian.com/technology/ng-inte...
1  https://www.theguardian.com/technology/ng-inte...
```

```

2 https://www.theguardian.com/technology/2023/oc...
2 https://www.theguardian.com/technology/2023/oc...
2 https://www.theguardian.com/technology/2023/oc...

```

```

                                apiUrl \
0 https://content.guardianapis.com/technology/20...
0 https://content.guardianapis.com/technology/20...
0 https://content.guardianapis.com/technology/20...
1 https://content.guardianapis.com/technology/ng...
1 https://content.guardianapis.com/technology/ng...
1 https://content.guardianapis.com/technology/ng...
1 https://content.guardianapis.com/technology/ng...
2 https://content.guardianapis.com/technology/20...
2 https://content.guardianapis.com/technology/20...
2 https://content.guardianapis.com/technology/20...

```

```

                                fields \
0 {'byline': 'Johana Bhuiyan', 'body': '<p>In <a...
0 {'byline': 'Johana Bhuiyan', 'body': '<p>In <a...
0 {'byline': 'Johana Bhuiyan', 'body': '<p>In <a...
1 {'byline': 'Hannah Devlin Science Corresponden...
1 {'byline': 'Hannah Devlin Science Corresponden...
1 {'byline': 'Hannah Devlin Science Corresponden...
1 {'byline': 'Hannah Devlin Science Corresponden...
2 {'byline': 'Kari Paul', 'body': '<p>Google is ...
2 {'byline': 'Kari Paul', 'body': '<p>Google is ...
2 {'byline': 'Kari Paul', 'body': '<p>Google is ...

```

```

                                tags isHosted pillarId \
0 {'id': 'technology/technology', 'type': 'keywo... False pillar/news
0 {'id': 'technology/artificialintelligenceai', ... False pillar/news
0 {'id': 'education/education', 'type': 'keyword... False pillar/news
1 {'id': 'technology/artificialintelligenceai', ... False pillar/news
1 {'id': 'technology/computing', 'type': 'keywor... False pillar/news
1 {'id': 'technology/technology', 'type': 'keywo... False pillar/news
1 {'id': 'uk/uk', 'type': 'keyword', 'sectionId'... False pillar/news
2 {'id': 'technology/alphabet', 'type': 'keyword... False pillar/news
2 {'id': 'technology/google', 'type': 'keyword',... False pillar/news
2 {'id': 'technology/technology', 'type': 'keywo... False pillar/news

```

```

pillarName                               byline \
0 News                                   Johana Bhuiyan
0 News                                   Johana Bhuiyan
0 News                                   Johana Bhuiyan
1 News Hannah Devlin Science Correspondent, Rich Cous...
1 News Hannah Devlin Science Correspondent, Rich Cous...
1 News Hannah Devlin Science Correspondent, Rich Cous...

```

1	News	Hannah Devlin	Science Correspondent, Rich Cous...
2	News		Kari Paul
2	News		Kari Paul
2	News		Kari Paul

	body	wordcount
0	<p>In <a href="https://www.theguardian.com/tec...	1585
0	<p>In <a href="https://www.theguardian.com/tec...	1585
0	<p>In <a href="https://www.theguardian.com/tec...	1585
1	<figure class="element element-atom element--i...	1741
1	<figure class="element element-atom element--i...	1741
1	<figure class="element element-atom element--i...	1741
1	<figure class="element element-atom element--i...	1741
2	<p>Google is doing well, but not well enough f...	554
2	<p>Google is doing well, but not well enough f...	554
2	<p>Google is doing well, but not well enough f...	554

If we check the length of the original `articles` dataframe against the new `tag_per_line` we can see that we have many more rows, one row per tag used in a story. We can also see the index values for our new dataframe are duplicated. This is because what was a single row, row 0 for example, is now three rows because the story had three tags. What was row 1 is now four rows, because the story at row 1 had four tags. These index values help us keep track of what rows ‘go together’ to make a single story, which we’ll need later.

```
[ ]: len(tag_per_line)
```

```
[ ]: 7666
```

```
[ ]: len(articles)
```

```
[ ]: 1173
```

Now our `tags` column is similar in structure to our `fields` column, which we unpacked earlier using `.json_normalize`. We can do the same again to generate a separate dataframe of `tag_data`. We also use a new method `.set_index()` to replace the default index that gets generated when `.json_normalize()` makes the new dataframe, with the index from `tag_per_line` which is keeping track of which rows go with which story.

```
[ ]: tag_data = pd.json_normalize(tag_per_line['tags'])
tag_data = tag_data.set_index(tag_per_line.index)
```

As we have index values keeping track of which row goes with which story, we could use those values to refer back to our original `articles` dataframe when we need additional information. However to simplify later tasks like printing out titles and urls, lets just copy the columns from `tag_per_line` into this new `tag_data` dataframe. The index lookup approach would be more space efficient but this isn’t an issue with data this size.

```
[ ]: tag_data['wordcount'] = tag_per_line['wordcount']
tag_data['article_title'] = tag_per_line['webTitle']
```

```
tag_data['article_url'] = tag_per_line['webUrl']
tag_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 7666 entries, 0 to 1997
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    7637 non-null   object
1   type                  7637 non-null   object
2   sectionId             7614 non-null   object
3   sectionName           7614 non-null   object
4   webTitle               7637 non-null   object
5   webUrl                 7637 non-null   object
6   apiUrl                7637 non-null   object
7   references             7637 non-null   object
8   description            391 non-null    object
9   activeSponsorships    2 non-null      object
10  wordcount              7666 non-null   int64
11  article_title          7666 non-null   object
12  article_url            7666 non-null   object
dtypes: int64(1), object(12)
memory usage: 838.5+ KB
```

We have quite a lot of columns in this dataset we've made, and probably only need a few. Let's just overwrite `tag_data` with a view of it that only includes the columns we need just to keep things simpler.

```
[ ]: tag_data = tag_data[['webTitle', 'article_title', 'article_url', 'wordcount']]
tag_data.head()
```

```
[ ]:
      webTitle \
0      Technology
0  Artificial intelligence (AI)
0      Education
1  Artificial intelligence (AI)
1      Computing

      article_title \
0  'Is this an appropriate use of AI or not?': te...
0  'Is this an appropriate use of AI or not?': te...
0  'Is this an appropriate use of AI or not?': te...
1      A day in the life of AI
1      A day in the life of AI

      article_url  wordcount
0  https://www.theguardian.com/technology/2023/oc...  1585
0  https://www.theguardian.com/technology/2023/oc...  1585
```

0	<a href="https://www.theguardian.com/technology/2023/oct...">https://www.theguardian.com/technology/2023/oct...</a>	1585
1	<a href="https://www.theguardian.com/technology/ng-inte...">https://www.theguardian.com/technology/ng-inte...</a>	1741
1	<a href="https://www.theguardian.com/technology/ng-inte...">https://www.theguardian.com/technology/ng-inte...</a>	1741

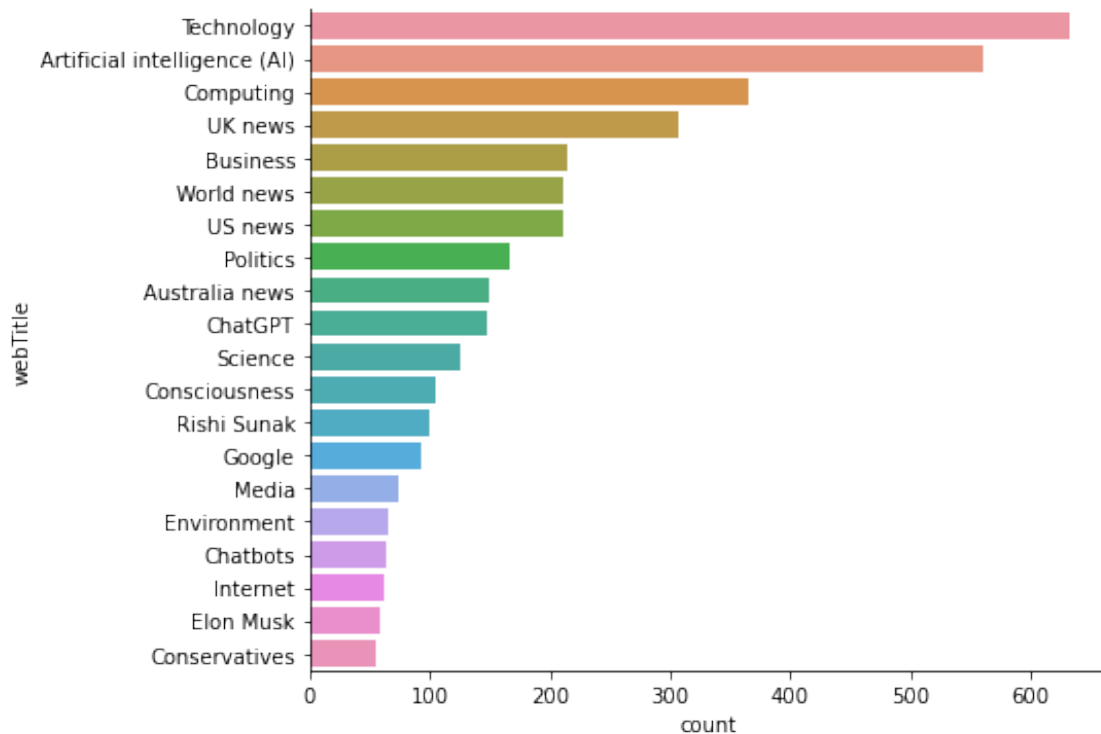
We can now check the count frequency of the different tags to get an overall picture like we did with sections.

```
[ ]: tag_counts = tag_data['webTitle'].value_counts().head(20)
      top_tags = tag_counts.index
      tag_counts
```

```
[ ]: Technology          633
      Artificial intelligence (AI)  560
      Computing          366
      UK news            307
      Business           214
      World news         211
      US news            211
      Politics           166
      Australia news     150
      ChatGPT            147
      Science            125
      Consciousness       105
      Rishi Sunak         100
      Google              92
      Media               73
      Environment         65
      Chatbots            63
      Internet            62
      Elon Musk           58
      Conservatives       55
      Name: webTitle, dtype: int64
```

```
[ ]: sns.catplot(data=tag_data, y='webTitle', kind='count', aspect=1.5,
      ↪order=top_tags)
```

```
[ ]: <seaborn.axisgrid.FacetGrid at 0x120a202e0>
```



## 1.8 Titles by Tag

Like before with sections, we can examine what stories are associated with each tag. The column names will be different but the mechanics are the same. N.B Below we use `.head()` to limit the number of results for demonstration purposes, but during analysis there is no reason you cannot remove it and to view all the results.

```
[ ]: TAG_OF_INTEREST = 'Elon Musk' # Just change this to switch tags

selected_data = tag_data[tag_data['webTitle'] == TAG_OF_INTEREST].head()

for index, row in selected_data.iterrows():
    print(row['article_title'])
    print(row['article_url'])
    print('****')
```

Five takeaways from UK's AI safety summit at Bletchley Park  
<https://www.theguardian.com/technology/2023/nov/02/five-takeaways-uk-ai-safety-summit-bletchley-park-rishi-sunak>

\*\*\*\*

Balancing the risks and rewards of AI will be key | Letters  
<https://www.theguardian.com/technology/2023/nov/06/balancing-the-risks-and-rewards-of-ai-will-be-key>

```

****
'Bletchley made me more optimistic': how experts reacted to AI summit
https://www.theguardian.com/technology/2023/nov/03/bletchley-made-me-more-
optimistic-how-experts-reacted-to-ai-summit
****
Sunak plays eager chatshow host as Musk discusses AI and politics
https://www.theguardian.com/politics/2023/nov/02/sunak-plays-eager-chatshow-
host-as-musk-discusses-ai-and-politics
****
Elon Musk unveils Grok, an AI chatbot with a 'rebellious streak'
https://www.theguardian.com/technology/2023/nov/05/elon-musk-unveils-grok-an-ai-
chatbot-with-a-rebellious-streak
****

```

We will use this data more later when examining wordcounts, and looking at tag correlation.

## 1.9 Word Counts

### 1.9.1 By Section

We defined `top_sections` earlier when we checked which sections had the highest number of stories. Here we'll use `.groupby` to get per section and per tag wordcounts. Word count is a good proxy for how much time was dedicated to a particular topic.

Total word count tells us the overall time dedicated to the topic related to each section or topic, whilst taking an average tells us how much space was given per story.

```

[ ]: section_wordcounts = articles.groupby('sectionName').agg(
    avg_wordcount=('wordcount', 'mean'),
    total_wordcount=('wordcount', 'sum')
).sort_values('total_wordcount', ascending=False).loc[top_sections]

section_wordcounts

```

```

[ ]:


```

	avg_wordcount	total_wordcount
Technology	854.544622	373436
Opinion	905.955224	182097
Australia news	2039.099099	226340
Business	1412.684211	134205
World news	1616.053763	150293
Politics	2216.072464	152909
US news	1786.863636	117933
Science	917.540541	33949
Environment	874.878788	28871
News	267.451613	8291

We can use box plots to see the distribution of these word counts. Remember we already filtered the `articles` data so it only included stories in top sections, however we include the filtering here to clarify that it is necessary before visualisation to reduce visual clutter.

```
[ ]: to_plot = articles[articles['sectionName'].isin(top_sections)]
sns.catplot(data=to_plot, y='sectionName', x='wordcount', kind='box', aspect=2)
```

```
[ ]: <seaborn.axisgrid.FacetGrid at 0x120defd30>
```



### 1.9.2 By Tag

For a similar summary, but by tag, we do the same, but we use the `tag_data` dataframe, and the `webTitle` column.

```
[ ]: tag_data.groupby('webTitle').agg(
    avg_wordcount=('wordcount', 'mean'),
    total_wordcount=('wordcount', 'sum')
).sort_values('total_wordcount', ascending=False).loc[top_tags]
```

```
[ ]:
```


	avg_wordcount	total_wordcount
Technology	908.581359	575132
Artificial intelligence (AI)	977.994643	547677
Computing	949.975410	347691
UK news	1153.117264	354007
Business	1263.654206	270422
World news	1004.488152	211947
US news	1245.033175	262702
Politics	1471.198795	244219
Australia news	1749.940000	262491
ChatGPT	849.571429	124887
Science	1017.304000	127163
Consciousness	962.542857	101067
Rishi Sunak	1849.910000	184991



Google	865.043478	79584
Media	900.136986	65710
Environment	1925.907692	125184
Chatbots	890.222222	56084
Internet	857.306452	53153
Elon Musk	900.931034	52254
Conservatives	2261.545455	124385

```
[ ]: to_plot = tag_data[tag_data['webTitle'].isin(top_tags)]
sns.catplot(data=to_plot, y='webTitle', x='wordcount', kind='box', aspect=2)
```

```
[ ]: <seaborn.axisgrid.FacetGrid at 0x120c8ceb0>
```



## 1.10 Tag Correlation

One analysis technique that is available to us is to examine the correlation of tags. What tags tend to co-occur in single stories, could this give us a sense of the themes or intersection of different topics?

Here we'll create a matrix of tag counts. In the first stage we use `.get_dummies` to reshape our column of tag names so that each possible tag is given its own column, and a value of 1 is entered if that tag is present in the row, otherwise 0.

(This may be a little confusing now but we're heading somewhere!)

```
[ ]: tag_matrix = pd.get_dummies(tag_data['webTitle'])
tag_matrix
```

```
[ ]:      3D  A-levels  AT&T  Abortion  Academics  Accountancy  Acting  Activism  \
0      0      0      0      0      0      0      0      0
0      0      0      0      0      0      0      0      0
0      0      0      0      0      0      0      0      0
1      0      0      0      0      0      0      0      0
1      0      0      0      0      0      0      0      0
... ..
1997  0      0      0      0      0      0      0      0
1997  0      0      0      0      0      0      0      0
1997  0      0      0      0      0      0      0      0
1997  0      0      0      0      0      0      0      0
1997  0      0      0      0      0      0      0      0
```

```
      Adam Bandt  Adobe  ...  YouTube  Young people  Youth unemployment  \
0      0      0      ...      0      0      0
0      0      0      ...      0      0      0
0      0      0      ...      0      0      0
1      0      0      ...      0      0      0
1      0      0      ...      0      0      0
...
1997      0      0      ...      0      0      0
1997      0      0      ...      0      0      0
1997      0      0      ...      0      0      0
1997      0      0      ...      0      0      0
1997      0      0      ...      0      0      0
```

```
      Yuval Noah Harari  Yvette Cooper  Zambia  Zoology  iOS  iPad  iPhone
0      0      0      0      0      0      0      0
0      0      0      0      0      0      0      0
0      0      0      0      0      0      0      0
1      0      0      0      0      0      0      0
1      0      0      0      0      0      0      0
...
1997      0      0      0      0      0      0      0
1997      0      0      0      0      0      0      0
1997      0      0      0      0      0      0      0
1997      0      0      0      0      0      0      0
1997      0      0      0      0      0      0      0
```

[7666 rows x 1032 columns]

Next we take our `tag_matrix`, use our list of `top_tags` to ensure only columns representing our selected top tags remain. We do this to aid visualisation later.

```
[ ]: tag_matrix = tag_matrix[top_tags].copy()
tag_matrix.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

Int64Index: 7666 entries, 0 to 1997

Data columns (total 20 columns):

#	Column	Non-Null Count	Dtype
0	Technology	7666 non-null	uint8
1	Artificial intelligence (AI)	7666 non-null	uint8
2	Computing	7666 non-null	uint8
3	UK news	7666 non-null	uint8
4	Business	7666 non-null	uint8
5	World news	7666 non-null	uint8
6	US news	7666 non-null	uint8
7	Politics	7666 non-null	uint8
8	Australia news	7666 non-null	uint8
9	ChatGPT	7666 non-null	uint8
10	Science	7666 non-null	uint8
11	Consciousness	7666 non-null	uint8
12	Rishi Sunak	7666 non-null	uint8
13	Google	7666 non-null	uint8
14	Media	7666 non-null	uint8
15	Environment	7666 non-null	uint8
16	Chatbots	7666 non-null	uint8
17	Internet	7666 non-null	uint8
18	Elon Musk	7666 non-null	uint8
19	Conservatives	7666 non-null	uint8

dtypes: uint8(20)

memory usage: 209.6 KB

At the moment our `tag_matrix` is one row per tag per story, meaning that for every row *only one* of those columns will have a number 1 in it to represent a tag is associated with that story. In order to understand if certain tags correlate, if they go together, we need to simplify so that one row represents a story, and each column shows either a 0 or a 1 depending on whether the tag is present in that story.

As stories can only use each tag once if we took all the rows for one single story, and for each column added the row values together, the result would be one row where 1 indicates if the tag is there or not because if it's not, we'd simply be adding together 0 for each row, resulting in 0. Using `groupby` we can grab each set of rows representing a single story, `.sum()` together the values in each column and then get one row back which provides this representation.

Let's demo this with a simplified example...

```
[ ]: toy_matrix = pd.read_csv('toy_matrix.csv')
toy_matrix
```

```
[ ]:   story  tag1  tag2  tag3
0     A     1     0     0
1     A     0     1     0
2     A     0     0     1
3     B     1     0     0
```

4	B	0	0	1
---	---	---	---	---

The `story` column just represents the id or title of the story, and then we have a column for each of three different tags. You'll see that each row only has one 1 in it because it is one row per tag. What we want is one row per story, so just two rows, one for story A, one for story B that puts the values spread across multiple rows into just one row.

You could probably do this in your head because really all we're saying is, for each story's subset of rows, if there is a 1 anywhere in the column, then the value is 1, otherwise it's 0. As we know a single story can only use a tag once, we can simplify this slightly complicated logic as just "grab all the rows for a story and for each column, add the values together".

```
[ ]: toy_matrix.groupby('story').sum()
```

```
[ ]:
      tag1  tag2  tag3
story
A         1     1     1
B         1     0     1
```

We can do the same with our actual `tag_matrix`. As we want to group on the index of the dataframe rather than a column we don't have a column name to pass `.groupby()` as usual. However we can tell it to group by "level 0". Pandas refers to indexes as levels and on a regular dataframe with just a single index, there is only one level, level 0.

```
[ ]: tag_matrix = tag_matrix.groupby(level=0).sum()
tag_matrix
```

```
[ ]:
      Technology  Artificial intelligence (AI)  Computing  UK news  Business  \
0              1              1              0              0              0
1              1              1              1              1              0
2              1              0              0              0              1
5              1              1              0              0              0
6              1              1              0              0              0
...           ...              ...              ...              ...              ...
1991            0              0              0              0              0
1992            0              0              0              0              0
1994            0              0              0              0              0
1996            0              0              0              0              0
1997            0              0              0              1              1

      World news  US news  Politics  Australia news  ChatGPT  Science  \
0              0        0          0              0          0          0
1              0        0          0              0          0          0
2              0        0          0              0          0          0
5              0        0          0              0          0          0
6              0        0          0              0          1          0
...           ...        ...          ...              ...          ...          ...
1991            1        0          0              0          0          0
1992            1        0          0              0          0          0
```

1994	0	1	0	0	0	0
1996	0	0	0	0	0	0
1997	0	0	1	0	0	0

	Consciousness	Rishi Sunak	Google	Media	Environment	Chatbots	\
0	0	0	0	0	0	0	
1	0	0	0	0	0	0	
2	0	0	1	0	0	0	
5	0	0	0	0	0	0	
6	0	0	0	0	0	1	
...	...	...	...	...	...	...	
1991	0	0	0	0	0	0	
1992	0	0	0	0	0	0	
1994	0	0	0	0	0	0	
1996	0	0	0	0	0	0	
1997	0	0	0	0	0	0	

	Internet	Elon Musk	Conservatives
0	0	0	0
1	0	0	0
2	0	0	0
5	0	0	0
6	0	0	0
...	...	...	...
1991	0	0	0
1992	0	0	0
1994	0	0	0
1996	0	0	0
1997	0	0	0

[1173 rows x 20 columns]

Finally we can get our correlation scores using `.corr`. This reshapes the data into a square, where both the rows and the columns represent tags, and the values represent the correlation between the two tags.

- 0 Represents no correlation
- 1 Represents the highest positive correlation, i.e. every story with tag **a** also includes tag **b**.
- A negative value indicates negative correlation, i.e. the presence of tag **a** means that the presence of tag **b** is less likely.

The ‘diagonal’ of the matrix will always equal 1 as the presence of tag **a** will always be correlated with the presence of tag **a**.

```
[ ]: correlations = tag_matrix.corr()
      correlations
```

```

[ ]:
Technology Artificial intelligence (AI) \
Technology 1.000000 0.574604
Artificial intelligence (AI) 0.574604 1.000000
Computing 0.533412 0.605133
UK news 0.016848 0.044402
Business 0.019890 -0.101796
World news 0.098574 0.050066
US news 0.000604 -0.052139
Politics -0.105899 0.013467
Australia news -0.266054 -0.161559
ChatGPT 0.266949 0.349627
Science 0.167821 0.188183
Consciousness 0.253656 0.328054
Rishi Sunak -0.061030 0.056591
Google 0.205829 0.171906
Media 0.103413 -0.006011
Environment -0.157571 -0.119598
Chatbots 0.182106 0.226543
Internet 0.157034 -0.019828
Elon Musk 0.163318 0.041806
Conservatives -0.191595 -0.074742

Computing UK news Business World news \
Technology 0.533412 0.016848 0.019890 0.098574
Artificial intelligence (AI) 0.605133 0.044402 -0.101796 0.050066
Computing 1.000000 0.034365 -0.055770 0.101387
UK news 0.034365 1.000000 0.099823 0.008971
Business -0.055770 0.099823 1.000000 0.008604
World news 0.101387 0.008971 0.008604 1.000000
US news -0.085447 -0.066770 0.077179 0.150511
Politics 0.001080 0.531680 -0.045873 -0.043685
Australia news -0.180739 -0.210569 -0.081274 -0.139455
ChatGPT 0.295293 -0.067208 -0.031944 0.043962
Science 0.224587 0.020464 -0.139635 0.060710
Consciousness 0.414030 0.003527 -0.108836 0.039749
Rishi Sunak 0.018436 0.380777 -0.041216 -0.007853
Google 0.159426 -0.051062 0.099730 0.020235
Media -0.036388 -0.040989 0.051625 -0.019582
Environment -0.106816 -0.076394 -0.046611 0.022389
Chatbots 0.206844 -0.038613 -0.014542 0.036109
Internet 0.038279 -0.001966 0.075426 -0.001514
Elon Musk 0.050099 -0.019501 -0.026135 -0.014672
Conservatives -0.105850 0.280779 -0.083416 -0.082877

US news Politics Australia news ChatGPT \
Technology 0.000604 -0.105899 -0.266054 0.266949
Artificial intelligence (AI) -0.052139 0.013467 -0.161559 0.349627

```

Computing	-0.085447	0.001080	-0.180739	0.295293
UK news	-0.066770	0.531680	-0.210569	-0.067208
Business	0.077179	-0.045873	-0.081274	-0.031944
World news	0.150511	-0.043685	-0.139455	0.043962
US news	1.000000	-0.094629	-0.166041	0.003738
Politics	-0.094629	1.000000	-0.155470	-0.109358
Australia news	-0.166041	-0.155470	1.000000	-0.067837
ChatGPT	0.003738	-0.109358	-0.067837	1.000000
Science	-0.081887	-0.013275	-0.114677	0.093756
Consciousness	-0.069098	-0.015930	-0.120065	0.224059
Rishi Sunak	-0.039646	0.603007	-0.116898	-0.106333
Google	0.036747	-0.100251	-0.064233	0.186482
Media	0.017169	-0.043847	-0.077510	-0.001581
Environment	-0.084329	-0.034195	0.018835	-0.069170
Chatbots	-0.032810	-0.075029	-0.091226	0.321008
Internet	-0.041194	-0.074050	-0.044820	0.048682
Elon Musk	0.046756	0.008935	-0.075559	-0.003190
Conservatives	-0.072377	0.523148	-0.084931	-0.083955

	Science	Consciousness	Rishi Sunak	Google \
Technology	0.167821	0.253656	-0.061030	0.205829
Artificial intelligence (AI)	0.188183	0.328054	0.056591	0.171906
Computing	0.224587	0.414030	0.018436	0.159426
UK news	0.020464	0.003527	0.380777	-0.051062
Business	-0.139635	-0.108836	-0.041216	0.099730
World news	0.060710	0.039749	-0.007853	0.020235
US news	-0.081887	-0.069098	-0.039646	0.036747
Politics	-0.013275	-0.015930	0.603007	-0.100251
Australia news	-0.114677	-0.120065	-0.116898	-0.064233
ChatGPT	0.093756	0.224059	-0.106333	0.186482
Science	1.000000	0.410658	-0.006437	0.012183
Consciousness	0.410658	1.000000	0.021907	0.108460
Rishi Sunak	-0.006437	0.021907	1.000000	-0.089060
Google	0.012183	0.108460	-0.089060	1.000000
Media	-0.065512	-0.031330	-0.078644	0.003603
Environment	0.060726	-0.062892	-0.033911	-0.042939
Chatbots	0.039921	0.150483	-0.072729	0.141484
Internet	-0.044147	-0.034032	-0.058473	0.101151
Elon Musk	-0.002283	0.080003	0.113432	-0.022657
Conservatives	-0.050016	-0.055420	0.495486	-0.064706

	Media	Environment	Chatbots	Internet \
Technology	0.103413	-0.157571	0.182106	0.157034
Artificial intelligence (AI)	-0.006011	-0.119598	0.226543	-0.019828
Computing	-0.036388	-0.106816	0.206844	0.038279
UK news	-0.040989	-0.076394	-0.038613	-0.001966
Business	0.051625	-0.046611	-0.014542	0.075426

World news	-0.019582	0.022389	0.036109	-0.001514
US news	0.017169	-0.084329	-0.032810	-0.041194
Politics	-0.043847	-0.034195	-0.075029	-0.074050
Australia news	-0.077510	0.018835	-0.091226	-0.044820
ChatGPT	-0.001581	-0.069170	0.321008	0.048682
Science	-0.065512	0.060726	0.039921	-0.044147
Consciousness	-0.031330	-0.062892	0.150483	-0.034032
Rishi Sunak	-0.078644	-0.033911	-0.072729	-0.058473
Google	0.003603	-0.042939	0.141484	0.101151
Media	1.000000	0.014728	-0.030066	0.175724
Environment	0.014728	1.000000	-0.057703	-0.057217
Chatbots	-0.030066	-0.057703	1.000000	0.214137
Internet	0.175724	-0.057217	0.214137	1.000000
Elon Musk	0.201686	-0.038053	-0.002007	0.121871
Conservatives	-0.023752	-0.018468	-0.052841	-0.034372

	Elon Musk	Conservatives
Technology	0.163318	-0.191595
Artificial intelligence (AI)	0.041806	-0.074742
Computing	0.050099	-0.105850
UK news	-0.019501	0.280779
Business	-0.026135	-0.083416
World news	-0.014672	-0.082877
US news	0.046756	-0.072377
Politics	0.008935	0.523148
Australia news	-0.075559	-0.084931
ChatGPT	-0.003190	-0.083955
Science	-0.002283	-0.050016
Consciousness	0.080003	-0.055420
Rishi Sunak	0.113432	0.495486
Google	-0.022657	-0.064706
Media	0.201686	-0.023752
Environment	-0.038053	-0.018468
Chatbots	-0.002007	-0.052841
Internet	0.121871	-0.034372
Elon Musk	1.000000	-0.013384
Conservatives	-0.013384	1.000000

We can check the correlations for a specific tag by accessing its column...

```
[ ]: correlations['ChatGPT'].sort_values(ascending=False)
```

```
[ ]: ChatGPT          1.000000
      Artificial intelligence (AI)  0.349627
      Chatbots         0.321008
      Computing        0.295293
      Technology       0.266949
      Consciousness    0.224059
```



Google	0.186482
Science	0.093756
Internet	0.048682
World news	0.043962
US news	0.003738
Media	-0.001581
Elon Musk	-0.003190
Business	-0.031944
UK news	-0.067208
Australia news	-0.067837
Environment	-0.069170
Conservatives	-0.083955
Rishi Sunak	-0.106333
Politics	-0.109358

Name: ChatGPT, dtype: float64


### 1.10.1 Tag Heatmap

We can also visualise these correlations using a heatmap. Using the `coolwarm` color scheme means colours run from deep blue to deep red. We set the `center` of the scale to 0 so that above zero, positive correlation, is a shade of red whilst below zero, negative correlation, is a shade of blue.

```
[ ]: plt.figure(figsize=(8,8))

sns.heatmap(correlations, cmap='coolwarm', center=0)
```

```
[ ]: <AxesSubplot:>
```



### 1.10.2 Advanced: Identifying multi-tag titles

What if you wanted to understand WHY two tags correlate. Perhaps ones that are unexpected. You will need to identify which stories have both tags using our `tag_matrix`, and then use the index values to look up the correct rows in the `articles`. We can then iterate over them and view title and url like before.

```
[ ]: TAG_1 = 'ChatGPT'
TAG_2 = 'Consciousness'
tag_filter = (tag_matrix[TAG_1] == 1) & (tag_matrix[TAG_2] == 1)

selected_story_index = tag_matrix[tag_filter].index
selected_story_index
```

```
[ ]: Int64Index([ 74, 102, 171, 181, 206, 207, 212, 225, 270, 299, 317,
                 342, 352, 424, 446, 460, 477, 478, 512, 524, 530, 563,
                 578, 670, 759, 768, 788, 793, 796, 804, 908, 945, 967,
                 975, 981, 1049, 1077, 1187],
                dtype='int64')
```

```
[ ]: selected_data = articles.loc[selected_story_index].head()

for index, row in selected_data.iterrows():
    print(row['webTitle'])
    print(row['webUrl'])
    print('****')
```

AI doomsday warnings a distraction from the danger it already poses, warns expert

<https://www.theguardian.com/technology/2023/oct/29/ai-doomsday-warnings-a-distraction-from-the-danger-it-already-poses-warns-expert>

\*\*\*\*

AI watch: from deepfakes to a rock star humanoid

<https://www.theguardian.com/technology/2023/jul/07/ai-watch-deepfakes-humanoid-robot-artificial-intelligence>

\*\*\*\*

Instead of banning AI, schools should use it to enhance learning | Letters

<https://www.theguardian.com/technology/2023/jul/09/instead-of-banning-ai-schools-should-use-it-to-enhance-learning>

\*\*\*\*

The professor's great fear about AI? That it becomes the boss from hell

<https://www.theguardian.com/technology/2023/aug/25/ai-artificial-intelligence-michael-wooldridge-christmas-royal-institution-lectures>

\*\*\*\*

The existential threat from AI - and from humans misusing it | Letters

<https://www.theguardian.com/technology/2023/jun/02/the-existential-threat-from-ai-and-from-humans-misusing-it>

\*\*\*\*

## 1.11 Summary

There will be many other ways in which this kind of data can be explored, depending on the kind of question you might have. However the above techniques give us a good overview of the data including the time period covered, the top topics, the type of content that has been collected (news, sport, opinion etc.) and allows us to get a sense of some correlations of the topics.

## 1.12 Exercises

Explore your own data set from the Guardian API. Use the techniques above to get a better sense of what you've collected.